

## [ Paper review 2 ]

---

# Bayesian Learning For Neural Networks ( Radford M. Neal, 1995)

---

## [ Contents ]

---

### 1. Introduction

1. Bayesian and frequentists views of learning
2. BNN (Bayesian Neural Networks)
3. MCMC (Markov Chain Monte Carlo)
4. Outline of the remainder of the thesis

### 2. Priors for Infinite Networks

1. Priors converging to Gaussian Processes
2. Priors converging to non-Gaussian stable distributions
3. Priors for networks with more than one hidden layer
4. Hierarchical Models

### 3. Monte Carlo Implementation

1. The hybrid Monte Carlo Algorithm
2. An implementation of Bayesian Neural Network Learning
3. A demonstration using hybrid Monte Carlo
4. Comparison of hybrid Monte Carlo with other methods
5. Variants of hybrid Monte Carlo

### 4. Evaluation of Neural Network Models

1. Network architectures, priors, and training procedures
2. Tests of the behavior of large networks
3. Tests of ARD (Automatic Relevance Determination) model
4. Tests of Bayesian Models on real data sets

### 5. Conclusion

## 1. Introduction

---

### 1-1. Bayesian and frequentists views of learning

---

- Estimate "Uncertainty"
- Not interested in the value of  $\theta$  it self,

but rather the "Value of some quantity that may be observed in the future" ( =  $x^{(n+1)}$  )

Likelihood, Posterior, Posterior Predictive distribution

- likelihood :  $L(\theta) = L(\theta | x^{(1)}, \dots, x^{(n)}) \propto P(x^{(1)}, \dots, x^{(n)} | \theta) = \prod_{i=1}^n P(x^{(i)} | \theta)$
- posterior :  $P(\theta | x^{(1)}, \dots, x^{(n)}) = \frac{P(x^{(1)}, \dots, x^{(n)} | \theta) P(\theta)}{P(x^{(1)}, \dots, x^{(n)})} \propto L(\theta | x^{(1)}, \dots, x^{(n)}) P(\theta)$
- posterior predictive :  $P(x^{(n+1)} | x^{(1)}, \dots, x^{(n)}) = \int P(x^{(n+1)} | \theta) P(\theta | x^{(1)}, \dots, x^{(n)}) d\theta$

Bayesian View

- prediction "by INTEGRATION" , "rather than MAXIMIZATION"
- avoids jumping to conclusion by considering not just the value of  $\theta$  ( that explains the data best), but also

other values of  $\theta$  that explain the data reasonably well

- maximum a posteriori probability (MAP)

MAP estimation is better characterized as a form of maximum PENALIZED likelihood estimation, with penalty being the "PRIOR" density

Hierarchical Models

- $\theta_k$  are independent given  $\gamma$

$$P(\theta) = P(\theta_1, \dots, \theta_p) = \int P(\gamma) \prod_{k=1}^p P(\theta_k | \gamma) d\gamma$$

- give the top-level hyperparameters prior distribution "that are very vague, or even improper" ( data is sufficiently informative! )

Learning Complex models

- From Bayesian view, adjusting the complexity of the model based on the amount of data "MAKES NO SENSE"
- if model and prior are correct for 1000 observation  $\rightarrow$  also correct for 10 observation
- Occam's Razor will be applied automatically!

## 1-2. BNN (Bayesian Neural Networks)

---

### 1-2-1. MLP ( Multi-layer Perceptron Networks )

neural networks are "NON-parametric"

- do have a parameter
- just more numerous, less interpretable than "parametric" models

MLP ( Multi-layer Perceptron Networks )

$$f_k(x) = b_k + \sum_j v_{jk} h_j(x)$$

$$h_j(x) = \tanh\left(a_j + \sum_i u_{ij} x_i\right)$$

- $u_{ij}$  : weight from "input unit  $i$ " to "hidden unit  $j$ "
- $v_{jk}$  : weight from "hidden unit  $j$ " to "output unit  $k$ "

( MLP can be used to define probabilistic models for Regression & Classification )

Regression

- Gaussian (  $y_k \sim N(f_k(x), \sigma_k^2)$  )
- $P(y | x) = \prod_k \frac{1}{\sqrt{2\pi}\sigma_k} \exp\left(-(f_k(x) - y_k)^2 / 2\sigma_k^2\right)$

Classification

- Softmax
- $P(y = k | x) = \exp(f_k(x)) / \sum_{k'} \exp(f_{k'}(x))$

Minimization of error function = MLE for the Gaussian noise model

If we use Cross Validation to find the appropriate weight penalty...

- problem 1 ) computationally expensive
- problem 2 ) if  $n$  networks find different local minima, for which different weight penalties are appropriate

Bayesian Goal : "to find the predictive distribution" in new case

- Posterior predictive distribution :

$$P(y^{(n+1)} | x^{(n+1)}, (x^{(1)}, y^{(1)}), \dots, (x^{(n)}, y^{(n)})) \\ = \int P(y^{(n+1)} | x^{(n+1)}, \theta) P(\theta | (x^{(1)}, y^{(1)}), \dots, (x^{(n)}, y^{(n)})) d\theta$$

- Likelihood :

$$L(\theta | (x^{(1)}, y^{(1)}), \dots, (x^{(n)}, y^{(n)})) = \prod_{i=1}^n P(y^{(i)} | x^{(i)}, \theta)$$

- single-point prediction :

$$\hat{y}_k^{(n+1)} = \int f_k(x^{(n+1)}, \theta) P(\theta | (x^{(1)}, y^{(1)}), \dots, (x^{(n)}, y^{(n)})) d\theta$$

(  $f_k$  : network output functions )

## 1-2-2. Selecting a network "model" & "prior"

give weights & biases a "Gaussian prior distribution"

advantages of hierarchical models

- hyperparameter : "variance" of the Gaussian prior for the weight & biases  
( allows the model to adapt to whatever degree of smoothness )

- role of hyperparameter (controlling the priors for weights ) = role of "weight decay"
- without the need for validation set!

## 1-2-3. The ARD(Automatic Relevance Determination) Model

### Problem

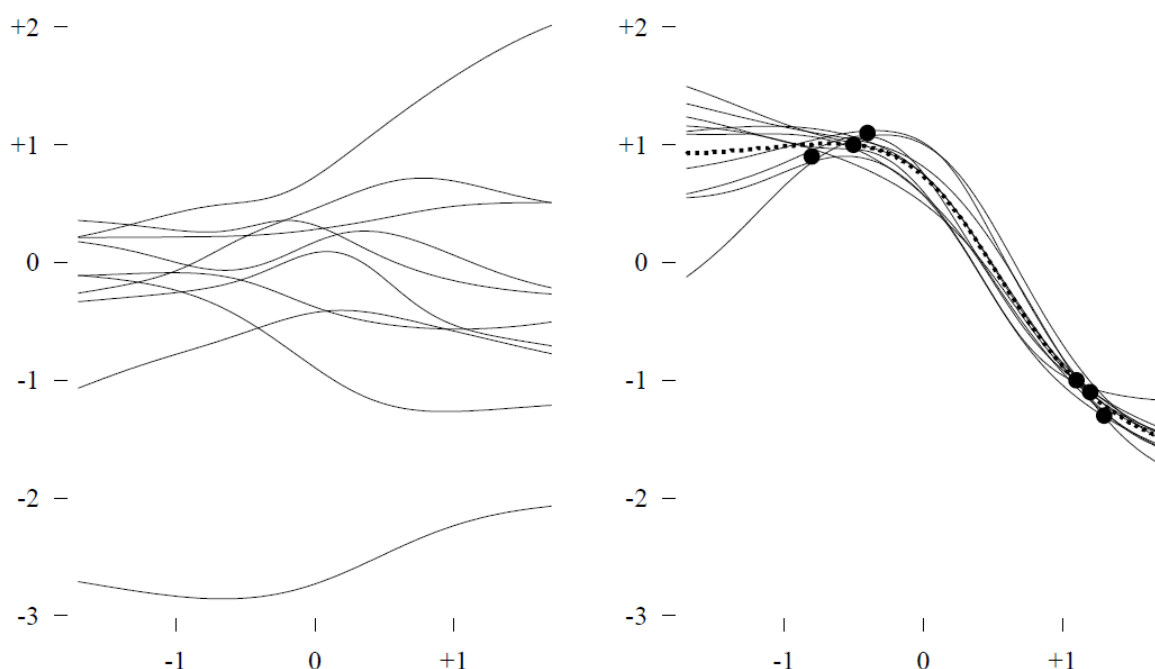
- Another dimension of complexity in NN = "number of input variables"  
( including more inputs → lead to poor generalization )
- Must limit the number of input variables, "based on our assessment of which attributes are most likely to be relevant"  
( So, which one is relevant ?)
- ARD : a model that can automatically determine the degree!

### ARD

- each input variable has associated with it a "hyperparameter" that controls the magnitudes of the weights  
( these hyperparameters are given some prior distribution )
- If hyperparameter associated with an input specifies small standard deviation for weights out of that input → weight will be small
- intended for use with a complex model ( with each input associated with many weights )  
( role of hyperparameter : introduce dependencies between these weights )
- will talk about more in Chapter 4

## 1-2-4. Illustration of Bayesian learning for Neural Networks

example) 1 input unit, 16 hidden units, 1 output units



(L) 10 networks, weight & bias from "Gaussian prior"

- easy to sample

(R) 10 networks, weight & bias from "Posterior distribution derived from prior & likelihood due to 6 data points"

- step 1 ) generate many networks from "prior"
- step 2 ) compute likelihood (for 6 points) of each 10 network
- accept each network ( from step 1 ) with a probability proportional to its likelihood ( step 2 )

Best way to guess the targets?

- use the AVERAGE of the network functions over the posterior distribution  
( how to average? → Monte Carlo Approximation )
- but Bayesian gives more than just a single-valued guess! also UNCERTAINTY  
( Uncertainty increases rapidly beyond the region where the training points are located )

## 1-2-5. Implementations based on Gaussian Approximation

"Finding the Predictive Distribution" = "Evaluating the Integral of Equation"

How to find the integral of equation?

- Ch 1) by Gaussian Approximation
- Ch 3) MCMC

Gaussian Approximation

- step 1) find (one or more) modes
- step 2) approximate posterior distribution  
( covariance matrix is chosen to match the second derivatives of the log posterior probability at the mode )
- step 3) If more than one mode → decide how much weight to give to each
- step 4) Approximate predictive distribution of equation  
( model that are linear in the vicinity of a mode, it is easy. If not, hard! )

Gaussian Approximation will fail to be good for LARGE networks

That's when we should use MCMC.

## 1-3. MCMC (Markov Chain Monte Carlo)

Make NO ASSUMPTION concerning the form of the distribution

deal with "Hybrid Monte Carlo" ( = Hamiltonian Monte Carlo )

### 1-3-1. Monte Carlo integration using Markov Chains

posterior :  $Q(\theta)$

then, expectation of  $a(\theta)$  :  $E[a] = \int a(\theta)Q(\theta)d\theta$

( let  $a(\theta) = f_k(x^{(n+1)}, \theta)$  to find the best guess for  $y_k^{(n+1)}$  )

How to find  $E[a] = \int a(\theta)Q(\theta)d\theta$ ?

- use Monte Carlo Approximation!

Monte Carlo Approximation

- $E[a] \approx \frac{1}{N} \sum_{t=1}^N a(\theta^{(t)})$
- $\theta^{(1)} \dots \theta^{(N)}$  are generated by distribution  $Q$  ( they are all independent )  
( even if  $Q$  is complicated .... would be possible to generate independent samples )
- gives UNBIASED estimate of  $E[a]$  ( even the samples are dependent! )
- still converge to true value (as long as the dependence is not great )

Invariance(=Stationary)

- $Q(\theta') = \int T(\theta' | \theta) Q(\theta)d\theta$
- stronger condition of "detailed balance"

Detailed balance

- $T(\theta' | \theta) Q(\theta) = T(\theta | \theta') Q(\theta')$
- "detailed balance"  $\rightarrow$  REVERSIBLE

"Ergodic"

- meaning that it has unique invariant(stationary) distribution ( = "Equilibrium distribution" )
- this distribution converges from ANY INITIAL state
- our goal : "construct a Markov chain which is ERGODIC"

To construct chain for complex problem :

- combine the transitions for simpler chains!
- sequence of it will also be invariant(stationary)

Check effect of dependencies of MC estimate? by auto-correlation !

Autocorrelation of  $a$  at lag  $s$  :  $\rho(s) = \frac{E\left[\left(a(\theta^{(t)}) - E[a]\right)\left(a(\theta^{(t-s)}) - E[a]\right)\right]}{\text{Var}[a]}$

### **1-3-2. Gibbs Sampling**

- skip

### **1-3-3. Metropolis Algorithm**

- skip

## **1-4. Outline of the remainder of the thesis**

Chapter 2)

- properties of prior distribution
- limit as the "number of hidden units  $\rightarrow \infty$  "
- goal ) reasonable priors (for such infinite networks) can be defined

Chapter 3)

- computational problems of BNN
- estimate using "hybrid MC"

Chapter 4)

- evaluate how good the predictions of BNNs are!
- evaluate the effectiveness of Hierarchical Models ( in particular, the ARD model )